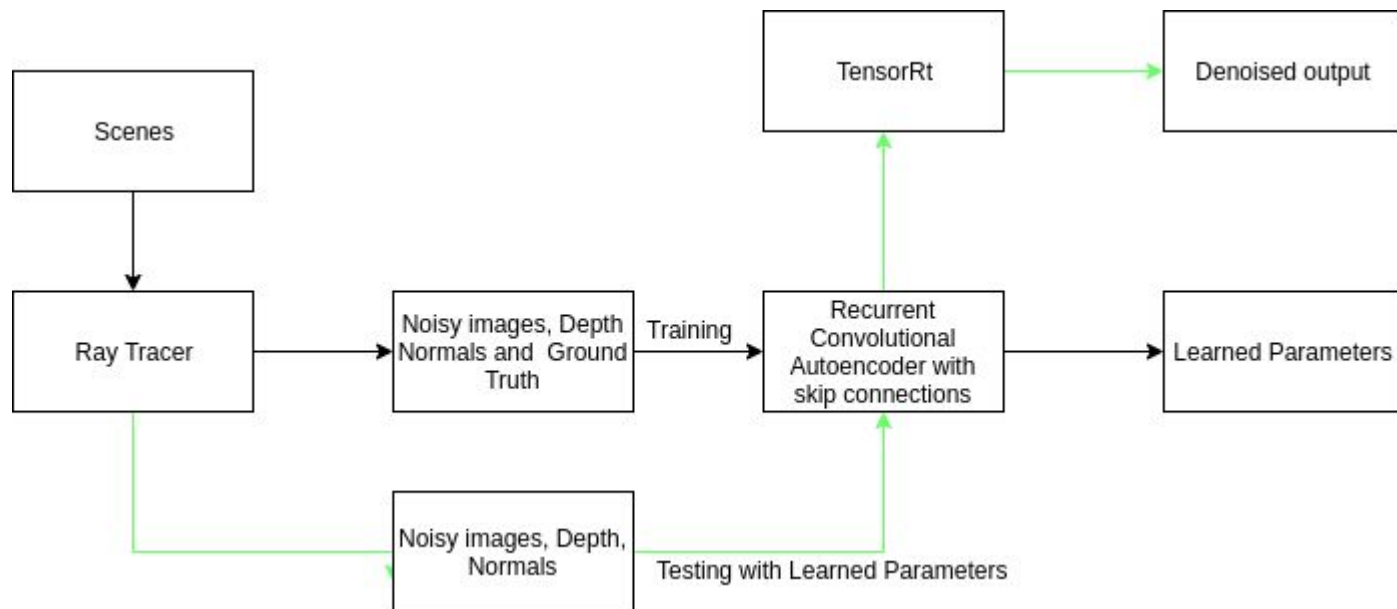




Another Image (AI) Denoiser


Dewang Sultania & Vaibhav Arcot

Workflow





Ray Tracer

- Features completed
 - Meshes
 - Dielectrics
 - Specular reflections
 - Parsing mtl files
 - Extracting normals, albedo and depth maps
- Features in progress
 - Texture loading
 - KD Tree



Data Generation

- Procedural Data Generation
 - Reflection, Refraction
 - Scale, Translation, Rotation
 - Color, Emittance
 - Camera eye
- Generated 80Gb of data
 - 30 scenes, 2503 images per scene



Architecture - Recurrent Convolutional Autoencoder with Skip Connections




Fig. 2. Architecture of our recurrent autoencoder. The input is 7 scalar values per pixel (noisy RGB, normal vector, depth, roughness). Each encoder stage applies a convolution and 2×2 max pooling. A decoder stage applies a 2×2 nearest neighbor upsampling, concatenates the per-pixel feature maps from connection (the spatial resolutions agree), and applies two sets of convolution and pooling. All convolutions have a 3×3 -pixel spatial support. On the right, we visualize the internal structure of the recurrent RCNN connections. I is the new input and h refers to the hidden, recurrent state that persists between animation frames.

- Convolutions
- Autoencoder
- Recurrent Layer
- Skip connections



Challenges

- Compute
 - Complicated architecture which requires a lot of GPU compute for training
- Data Storage:
 - 30 scenes with 2503 images each are 80 GB worth of data
 - Resizing and storing might be a solution, but that happens at the expense of performance.



Next Steps

- Benchmark PyTorch model
- Implement the neural network in cuDNN.
- Generate more scenes with textures.
- Speed up with TensorRt